

```

Apr 18, 25 4:42      EX2.asciidoc      Page 1/3

**In[1]:**
[source, ipython3]
----
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
----

**In[5]:**
[source, ipython3]
----
# Load data
df = pd.read_csv('encoded_car_data.csv')
print(df.head())
# Select features & target
X = df[['engine size', 'horsepower', 'citympg', 'highwaympg']]
y = df['price']
----

**Out[5]:**
----
   diesel  gas  std  turbo  convertible  hardtop  hatchback  sedan  wagon  \
0      0.0  1.0  1.0   0.0           1.0     0.0         0.0   0.0   0.0
1      0.0  1.0  1.0   0.0           1.0     0.0         0.0   0.0   0.0
2      0.0  1.0  1.0   0.0           0.0     0.0         1.0   0.0   0.0
3      0.0  1.0  1.0   0.0           0.0     0.0         0.0   1.0   0.0
4      0.0  1.0  1.0   0.0           0.0     0.0         0.0   1.0   0.0

   4wd  ...  wheelbase  curbweight  enginesize  boreratio  horsepower  \
0  0.0  ...      88.6      2548.0      130.0      3.47      111.0
1  0.0  ...      88.6      2548.0      130.0      3.47      111.0
2  0.0  ...      94.5      2823.0      152.0      2.68      154.0
3  0.0  ...      99.8      2337.0      109.0      3.19      102.0
4  1.0  ...      99.4      2824.0      136.0      3.19      115.0

   carlength  carwidth  citympg  highwaympg  price
0      168.8      64.1     21.0      27.0  13495.0
1      168.8      64.1     21.0      27.0  16500.0
2      171.2      65.5     19.0      26.0  16500.0
3      176.6      66.2     24.0      30.0  13950.0
4      176.6      66.4     18.0      22.0  17450.0

[5 rows x 36 columns]
----

**In[6]:**
[source, ipython3]
----
# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
state=42)
----

**In[7]:**
[source, ipython3]
----
# 1. Linear Regression (with scaling)
linear_model = Pipeline([
    ('scaler', StandardScaler()),
    ('model', LinearRegression())
])

```

```

Apr 18, 25 4:42      EX2.asciidoc      Page 2/3

linear_model.fit(X_train, y_train)
y_pred_linear = linear_model.predict(X_test)
----

**In[14]:**
[source, ipython3]
----
#2. Polynomial Regression (degree=2)
poly_model = Pipeline([
    ('poly', PolynomialFeatures(degree=2)),
    ('scaler', StandardScaler()),
    ('model', LinearRegression())
])
poly_model.fit(X_train, y_train)
y_pred_poly = poly_model.predict(X_test)
----

**In[15]:**
[source, ipython3]
----
# Evaluate Models
print("Linear Regression:")
print(f"MSE: {mean_squared_error(y_test, y_pred_linear):.2f}")
print(f"R2: {r2_score(y_test, y_pred_linear):.2f}")

print("\nPolynomial Regression:")
print(f"MSE: {mean_squared_error(y_test, y_pred_poly):.2f}")
print(f"R2: {r2_score(y_test, y_pred_poly):.2f}")
----

**Out[15]:**
----
Linear Regression:
MSE: 16471505.90
R2: 0.79

Polynomial Regression:
MSE: 15247661.89
R2: 0.81
----

**In[16]:**
[source, ipython3]
----
# Plot actual vs predicted

plt.figure(figsize=(10, 5))
plt.scatter(y_test, y_pred_linear, label='Linear', alpha=0.6)
plt.scatter(y_test, y_pred_poly, label='Polynomial (degree=2)', alpha=0.6)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', label='Perfect Predictio
n')
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Linear vs Polynomial Regression")
plt.legend()
plt.show()
----

**Out[16]:**
----
! [png] (output_6_0.png)
----

```

Apr 18, 25 4:42	EX2.asciidoc	Page 3/3
<pre data-bbox="73 188 284 284">**In[]:**+ [source, ipython3] ----- -----</pre>		